

From Artificial Neural Networks to Deep Learning


Dr. Konda Reddy Mopuri
kmopuri@iittp.ac.in
Dept. of CSE, IIT Tirupati
Aug-Dec 2021

What is DL?

What is DL?

- Subset of ML that is essentially neural networks with more layers
- Crude attempt to imitate the human brain in learning

What is DL?



Classical ML vs. DL

- Classical ML: Handcrafted features + learnable model
- Need strong domain expertise

Classical ML vs. DL

- Classical ML: Handcrafted features + learnable model
- Need strong domain expertise




Figure credits: Jay Shaw & Quora

Classical ML vs. DL

- Deep Learning: Deep stack of parameterized processing
- End-to-End learning

Classical ML vs. DL

- Deep Learning: Deep stack of parameterized processing
- End-to-End learning




Figure credits: Jay Shaw & Quora

Classical ML vs. DL

- ANNs predate some of the classical ML techniques
- We are now dealing with a new generation ANNs

Neuron

- About 100 billion neurons in human brain




Figure credits: Wikipedia

History of Neural Networks

- ① McCulloch Pitts neuron (1943) - Threshold Logic Unit

History of Neural Networks

- ① McCulloch Pitts neuron (1943) - Threshold Logic Unit
- ② Donald Hebb (1949) - Hebbian Learning Principle

History of Neural Networks

- ① McCulloch Pitts neuron (1943) - Threshold Logic Unit
- ② Donald Hebb (1949) - Hebbian Learning Principle
- ③ Marvin Minsky (1951) - created the first ANN (Hebbian Learning, 40 neurons)

History of Neural Networks

- ① McCulloch Pitts neuron (1943) - Threshold Logic Unit
- ② Donald Hebb (1949) - Hebbian Learning Principle
- ③ Marvin Minsky (1951) - created the first ANN (Hebbian Learning, 40 neurons)
- ④ Frank Rosenblatt (1958) - created perceptron to classify 20X20 images

History of Neural Networks

- ① McCulloch Pitts neuron (1943) - Threshold Logic Unit
- ② Donald Hebb (1949) - Hebbian Learning Principle
- ③ Marvin Minsky (1951) - created the first ANN (Hebbian Learning, 40 neurons)
- ④ Frank Rosenblatt (1958) - created perceptron to classify 20X20 images
- ⑤ David H Hubel and Torsten Wiesel (1959) demonstrated orientation selectivity and columnar organization in cat's visual cortex

Backpropagation

- Paul Werbos (1982) proposed back-propagation for ANNs

History (contd.)

- ① Neocognitron -Network for TC problem

History (contd.)

- ① Neocognitron -Network for TC problem
- ② LeNet family (Lecun et al. 1989)

History (contd.)

- ① Neocognitron -Network for TC problem
- ② LeNet family (Lecun et al. 1989)
- ③ AlexNet (Alex krizhevsky et al. 2012)

History (contd.)

- ① Neocognitron -Network for TC problem
- ② LeNet family (Lecun et al. 1989)
- ③ AlexNet (Alex krizhevsky et al. 2012)
- ④ GoogLeNet and ResNet, etc.

History (contd.)

- ① Neocognitron -Network for TC problem
- ② LeNet family (Lecun et al. 1989)
- ③ AlexNet (Alex krizhevsky et al. 2012)
- ④ GoogLeNet and ResNet, etc.
- ⑤ Transformers

Deep Learning

- ① Natural generalization to ANNs - Doesn't differ much from the 90s NNs


Deep Learning

- ① Natural generalization to ANNs - Doesn't differ much from the 90s NNs
- ② Computational graph of tensor operations that take advantage of
 - Chain rule (back-propagation)
 - SGD
 - GPUs
 - Huge datasets
 - Convolutions, etc.

Deep Learning

- This generalization enables us to build complex networks that work with Images, text, speech and sequences and train end-to-end

ILSVRC Error



Gershgorn 2017

What makes it work now?

What makes it work now?

- ① Huge research and progress in ML

What makes it work now?

- ① Huge research and progress in ML
- ② Hardware developments - CPUs/GPUs/Storage technologies

What makes it work now?

- ① Huge research and progress in ML
- ② Hardware developments - CPUs/GPUs/Storage technologies
- ③ Piles of data over the Internet

What makes it work now?

- ① Huge research and progress in ML
- ② Hardware developments - CPUs/GPUs/Storage technologies
- ③ Piles of data over the Internet
- ④ Collaborative development (open source tools and forums for sharing/discussions, etc)

What makes it work now?

- ① Huge research and progress in ML
- ② Hardware developments - CPUs/GPUs/Storage technologies
- ③ Piles of data over the Internet
- ④ Collaborative development (open source tools and forums for sharing/discussions, etc)
- ⑤ Collective efforts from large institutions/corporations

What makes it work now?

- ① Huge research and progress in ML
- ② Hardware developments - CPUs/GPUs/Storage technologies
- ③ Piles of data over the Internet
- ④ Collaborative development (open source tools and forums for sharing/discussions, etc)
- ⑤ Collective efforts from large institutions/corporations
- ⑥ ...

What makes it work now?

- We have been doing a lot of ML already
 - Taxonomy of ML concepts: Classification, regression, generative models, clustering, etc.
 - Rich statistical formalizations: Bayesian estimation, PAC, etc.
 - Understood fundamentals: Bias-Variance, VC dimension, etc.
 - Good understanding of optimization
 - Efficient large-scale algorithms

Deep Learning - practical perspective

- ① Doesn't require a deep mathematical grasp

Deep Learning - practical perspective

- ① Doesn't require a deep mathematical grasp
- ② Makes the design of large models a system/software development task

Deep Learning - practical perspective

- ① Doesn't require a deep mathematical grasp
- ② Makes the design of large models a system/software development task
- ③ Leverages modern hardware

Deep Learning - practical perspective

- ① Doesn't require a deep mathematical grasp
- ② Makes the design of large models a system/software development task
- ③ Leverages modern hardware
- ④ Doesn't seem to plateau with more data

Deep Learning - practical perspective

- ① Doesn't require a deep mathematical grasp
- ② Makes the design of large models a system/software development task
- ③ Leverages modern hardware
- ④ Doesn't seem to plateau with more data
- ⑤ Makes the trained models a commodity

Implementation

	Language(s)	License	Main backer
PyTorch	Python, C++	BSD	Facebook
TensorFlow	Python, C++	Apache	Google
JAX	Python	Apache	Google
MXNet	Python, C++, R, Scala	Apache	Amazon
CNTK	Python, C++	MIT	Microsoft
Torch	Lua	BSD	Facebook
Theano	Python	BSD	U. of Montreal
Caffe	C++	BSD 2 clauses	U. of CA, Berkeley

We use PyTorch



<http://pytorch.org>